

SYNTHETIC DATA SET CREATION

fds-blog

January 22, 2024

```
[260]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[261]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import sklearn
import warnings
import random
warnings.filterwarnings('ignore')
```

```
[262]: pip install faker
```

Requirement already satisfied: faker in /usr/local/lib/python3.10/dist-packages (22.4.0)

Requirement already satisfied: python-dateutil>=2.4 in /usr/local/lib/python3.10/dist-packages (from faker) (2.8.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.4->faker) (1.16.0)

1 Creating Synthetic Dataset Generation using Faker

```
[263]: from faker import Faker
np.random.seed(32)
fake = Faker()
Faker.seed(32)

# Number of rows in the dataset
start_index = 1
num_rows = 1500
# Generate synthetic data
battery_values = [3000, 3500, 4000, 4500, 5000, 5500, 6000]

data = {
```

```

    'Unnamed: 0': list(range(start_index, start_index + num_rows)),
    'Company': [fake.random_element(['Apple', 'Samsung', 'Google', 'Huawei',
    ↪ 'OnePlus', 'Sony', 'LG', 'Motorola', 'Xiaomi', 'Nokia', 'Oppo', 'Vivo',
    ↪ 'Realme']) for _ in range(num_rows)],
    'Weight(gm)': [fake.pyfloat(right_digits=2, positive=True, min_value=100,
    ↪ max_value=250) for _ in range(num_rows)],
    'PPI': [fake.random_int(min=100, max=800) for _ in range(num_rows)],
    'CPU_core': [fake.random_int(min=1, max=8) for _ in range(num_rows)],
    'CPU_freq': [fake.pyfloat(left_digits=1, right_digits=1, positive=True,
    ↪ min_value=1, max_value=2.7) for _ in range(num_rows)],
    'Dual_sim': [fake.random_element(['Yes', 'No']) for _ in range(num_rows)],
    'Internal_mem(GB)': [fake.random_element(['16GB', '32GB',
    ↪ '64GB', '128GB', '256GB']) for _ in range(num_rows)],
    'RAM': [fake.random_int(min=4, max=16) for _ in range(num_rows)],
    'RearCam': [fake.random_int(min=20, max=108) for _ in range(num_rows)],
    'Front_Cam': [fake.random_int(min=10, max=32) for _ in range(num_rows)],
    'Gen_5G': [fake.random_element(['Yes', 'No']) for _ in range(num_rows)],
    'Battery': [fake.random_element(battery_values) for _ in range(num_rows)],
    'Thickness': [fake.random_int(min=5, max=18) for _ in range(num_rows)],
    # 'Price': [fake.random_int(min=5000, max=100000) for _ in range(num_rows)]
}

df = pd.DataFrame(data)
df.head()

```

```

[263]:
   Unnamed: 0  Company  Weight(gm)  PPI  CPU_core  CPU_freq  Dual_sim  \
0           1  Samsung    180.32   312         1        1.5        No
1           2   Huawei    159.12   362         1        1.6        No
2           3   Google    159.29   241         1        1.0        No
3           4  OnePlus    214.38   555         4        1.2        Yes
4           5    Vivo     102.43   607         8        1.8        Yes

   Internal_mem(GB)  RAM  RearCam  Front_Cam  Gen_5G  Battery  Thickness
0             16GB    8       77         31    Yes     5000         17
1             64GB    8       51         13    No     3000         12
2             16GB   11       84         23    No     5000         11
3             32GB    7       91         10    No     6000          8
4            128GB   12       71         27    No     4500         11

```

2 Removing unwanted rows and unwanted data

```

[264]: df.drop(columns=['Unnamed: 0'], inplace=True)
df["Internal_mem(GB)"] = df["Internal_mem(GB)"].str.replace('GB', '').
    ↪ astype(int)

```

```
[265]: df_copy = df.copy()
df_copy.head()
```

```
[265]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	Internal_mem(GB)	\
0	Samsung	180.32	312	1	1.5	No	16	
1	Huawei	159.12	362	1	1.6	No	64	
2	Google	159.29	241	1	1.0	No	16	
3	OnePlus	214.38	555	4	1.2	Yes	32	
4	Vivo	102.43	607	8	1.8	Yes	128	

	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	
0	8	77	31	Yes	5000	17	
1	8	51	13	No	3000	12	
2	11	84	23	No	5000	11	
3	7	91	10	No	6000	8	
4	12	71	27	No	4500	11	

3 Performing label encoding (converting catogirical data to numerical data)

```
[266]: from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
df_copy['Gen_5G'] = label_encoder.fit_transform(df_copy['Gen_5G']).astype(int)
df_copy['Dual_sim'] = label_encoder.fit_transform(df_copy['Dual_sim']).
↳astype(int)
```

```
[267]: df_copy.describe()
```

```
[267]:
```

	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.000000	
mean	176.520353	449.583333	4.427333	1.451667	0.518667	
std	43.239296	200.174228	2.345006	0.287249	0.499818	
min	100.120000	100.000000	1.000000	1.000000	0.000000	
25%	139.862500	272.750000	2.000000	1.200000	0.000000	
50%	175.975000	451.000000	4.000000	1.500000	1.000000	
75%	214.347500	624.000000	6.000000	1.700000	1.000000	
max	249.940000	799.000000	8.000000	1.900000	1.000000	

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	\
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.000000	
mean	98.570667	10.054667	63.491333	21.166667	0.491333	
std	86.612660	3.786983	25.918573	6.680767	0.500092	
min	16.000000	4.000000	20.000000	10.000000	0.000000	
25%	32.000000	7.000000	41.000000	15.000000	0.000000	
50%	64.000000	10.000000	63.000000	21.000000	0.000000	
75%	128.000000	14.000000	86.000000	27.000000	1.000000	

max	256.000000	16.000000	108.000000	32.000000	1.000000
-----	------------	-----------	------------	-----------	----------

	Battery	Thickness
count	1500.000000	1500.000000
mean	4525.333333	11.478000
std	1006.992674	4.028938
min	3000.000000	5.000000
25%	3500.000000	8.000000
50%	4500.000000	11.000000
75%	5500.000000	15.000000
max	6000.000000	18.000000

4 Generating Price based on the features weights(How much the feature contributes in price)

```
[268]: #company wegihts
user_defined_values = {'Apple': 250, 'Samsung': 160, 'Google': 85, 'Huawei': 100, 'OnePlus': 180, 'Sony': 90, 'LG': 70,
                        'Motorola': 60, 'Xiaomi': 45, 'Nokia': 50, 'Oppo': 20, 'Vivo': 40, 'Realme': 30,}
df_copy['Company'] = df_copy['Company'].map(user_defined_values)
```

5 Maximum Absolute scaler

```
[269]: max_vals = np.max(np.abs(df_copy))
max_vals
```

```
[269]: Company          250.00
Weight(gm)          249.94
PPI                  799.00
CPU_core              8.00
CPU_freq              1.90
Dual_sim              1.00
Internal_mem(GB)      256.00
RAM                   16.00
RearCam              108.00
Front_Cam            32.00
Gen_5G                1.00
Battery              6000.00
Thickness             18.00
dtype: float64
```

```
[270]: r=(df_copy) / max_vals
r.head()
```

```
[270]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	0.64	0.721453	0.390488	0.125	0.789474	0.0	
1	0.40	0.636633	0.453066	0.125	0.842105	0.0	
2	0.34	0.637313	0.301627	0.125	0.526316	0.0	
3	0.72	0.857726	0.694618	0.500	0.631579	1.0	
4	0.16	0.409818	0.759700	1.000	0.947368	1.0	

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness
0	0.0625	0.5000	0.712963	0.96875	1.0	0.833333	0.944444
1	0.2500	0.5000	0.472222	0.40625	0.0	0.500000	0.666667
2	0.0625	0.6875	0.777778	0.71875	0.0	0.833333	0.611111
3	0.1250	0.4375	0.842593	0.31250	0.0	1.000000	0.444444
4	0.5000	0.7500	0.657407	0.84375	0.0	0.750000	0.611111

```
[271]: def calculate_price(data):
# weight_vector = [17,1,7,8,7,4,14,15,7,6,3,13,2]
weight_vector = [0.15, 0.04, 0.05, 0.11, 0.08, 0.04, 0.14, 0.17, 0.06, 0.02, 0.04, 0.09, 0.01]
x = pd.DataFrame(np.array(weight_vector * data))
x = x*1000
x.columns = df.columns
return x
df_new = calculate_price(r)
df_new.head()
```

```
[271]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	96.0	28.858126	19.524406	13.75	63.157895	0.0	
1	60.0	25.465312	22.653317	13.75	67.368421	0.0	
2	51.0	25.492518	15.081352	13.75	42.105263	0.0	
3	108.0	34.309034	34.730914	55.00	50.526316	40.0	
4	24.0	16.392734	37.984981	110.00	75.789474	40.0	

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness
0	8.75	85.000	42.777778	19.375	40.0	75.0	9.444444
1	35.00	85.000	28.333333	8.125	0.0	45.0	6.666667
2	8.75	116.875	46.666667	14.375	0.0	75.0	6.111111
3	17.50	74.375	50.555556	6.250	0.0	90.0	4.444444
4	70.00	127.500	39.444444	16.875	0.0	67.5	6.111111

```
[272]: df_new['Price'] = df_new.sum(axis=1)
df_new['Price'] = df_new['Price'].round(-1)
df_new
```

```
[272]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	96.0	28.858126	19.524406	13.75	63.157895	0.0	
1	60.0	25.465312	22.653317	13.75	67.368421	0.0	
2	51.0	25.492518	15.081352	13.75	42.105263	0.0	

3	108.0	34.309034	34.730914	55.00	50.526316	40.0
4	24.0	16.392734	37.984981	110.00	75.789474	40.0
...
1495	42.0	34.566696	39.236546	41.25	67.368421	0.0
1496	150.0	24.767544	28.848561	27.50	54.736842	40.0
1497	54.0	33.880131	29.161452	55.00	58.947368	0.0
1498	42.0	16.115868	46.433041	27.50	71.578947	40.0
1499	27.0	35.048412	42.365457	55.00	63.157895	0.0

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	\
0	8.75	85.000	42.777778	19.375	40.0	75.0	
1	35.00	85.000	28.333333	8.125	0.0	45.0	
2	8.75	116.875	46.666667	14.375	0.0	75.0	
3	17.50	74.375	50.555556	6.250	0.0	90.0	
4	70.00	127.500	39.444444	16.875	0.0	67.5	
...	
1495	8.75	116.875	30.000000	6.875	40.0	67.5	
1496	17.50	42.500	35.000000	15.625	0.0	75.0	
1497	140.00	74.375	42.777778	15.625	0.0	90.0	
1498	17.50	63.750	23.888889	19.375	0.0	45.0	
1499	140.00	53.125	23.888889	13.750	40.0	52.5	

	Thickness	Price
0	9.444444	500.0
1	6.666667	400.0
2	6.111111	420.0
3	4.444444	570.0
4	6.111111	630.0
...
1495	8.333333	500.0
1496	3.333333	510.0
1497	7.222222	600.0
1498	2.777778	420.0
1499	3.333333	550.0

[1500 rows x 14 columns]

```
[273]: # df_new['Price'].value_counts()
```

```
[274]: df_merged = pd.merge(df, df_new[['Price']], left_index=True, right_index=True)
```

```
[275]: df_merged.describe()
```

```
[275]:
```

	Weight(gm)	PPI	CPU_core	CPU_freq	Internal_mem(GB)	\
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.000000	
mean	176.520353	449.583333	4.427333	1.451667	98.570667	
std	43.239296	200.174228	2.345006	0.287249	86.612660	

min	100.120000	100.000000	1.000000	1.000000	16.000000
25%	139.862500	272.750000	2.000000	1.200000	32.000000
50%	175.975000	451.000000	4.000000	1.500000	64.000000
75%	214.347500	624.000000	6.000000	1.700000	128.000000
max	249.940000	799.000000	8.000000	1.900000	256.000000

	RAM	RearCam	Front_Cam	Battery	Thickness \
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.000000
mean	10.054667	63.491333	21.166667	4525.333333	11.478000
std	3.786983	25.918573	6.680767	1006.992674	4.028938
min	4.000000	20.000000	10.000000	3000.000000	5.000000
25%	7.000000	41.000000	15.000000	3500.000000	8.000000
50%	10.000000	63.000000	21.000000	4500.000000	11.000000
75%	14.000000	86.000000	27.000000	5500.000000	15.000000
max	16.000000	108.000000	32.000000	6000.000000	18.000000

	Price
count	1500.000000
mean	558.093333
std	91.443640
min	310.000000
25%	490.000000
50%	550.000000
75%	620.000000
max	860.000000

```
[276]: # Base_Price = {'Apple': 30000, 'Samsung': 5000, 'Google': 5000, 'Huawei': 4500, 'OnePlus': 8000, 'Sony': 3000, 'LG': 5000,
#           'Motorola': 6000, 'Xiaomi': 4500, 'Nokia': 6000, 'Oppo': 5000, 'Vivo': 5000, 'Realme': 5500,}
Base_Price = {'Apple': 30000, 'Samsung': 16000, 'Google': 20000, 'Huawei': 15000, 'OnePlus': 17000, 'Sony': 9000, 'LG': 6000,
              'Motorola': 4500, 'Xiaomi': 6500, 'Nokia': 6000, 'Oppo': 5000, 'Vivo': 7000, 'Realme': 8000,}

base_price_cpucore = {1: 500, 2: 1000, 3: 1500, 4: 2000, 5: 2500, 6: 3000, 7: 3500, 8: 4000,}
base_price_dual_sim = {'Yes': 3000, 'No': 0,}
base_price_Internal_mem = {16: 2000, 32: 3000, 64: 4000, 128: 5000, 256: 6000,}
base_price_RAM = {4: 500, 5: 1000, 6: 1500, 7: 2000, 8: 2500, 9: 3000, 10: 3000, 11: 4000, 12: 5000, 13: 6000, 14: 6000, 15: 6000, 16: 7000,}
base_price_Gen_5G = {'Yes': 4000, 'No': 0,}
base_price_battery = {3000: 1000, 3500: 1500, 4000: 2000, 4500: 2500, 5000: 3000, 5500: 4000, 6000: 5500,}

# Map user-defined values to the 'Company' column
```

```

df_merged['Base_Price'] = df_merged['Company'].map(Base_Price)
df_merged['base_price_cpucore'] = df_merged['CPU_core'].map(base_price_cpucore)
df_merged['base_price_dual_sim'] = df_merged['Dual_sim'].
    ↳map(base_price_dual_sim)
df_merged['base_price_Internal_mem'] = df_merged['Internal_mem(GB)'].
    ↳map(base_price_Internal_mem)
df_merged['base_price_RAM'] = df_merged['RAM'].map(base_price_RAM)
df_merged['base_price_Gen_5G'] = df_merged['Gen_5G'].map(base_price_Gen_5G)
df_merged['base_price_battery'] = df_merged['Battery'].map(base_price_battery)

for i in range(3):
    df_merged['Price'] += df_merged['Base_Price']

for i in range(1):
    df_merged['Price'] += df_merged['base_price_cpucore']

for i in range(1):
    df_merged['Price'] += df_merged['base_price_dual_sim']

for i in range(1):
    df_merged['Price'] += df_merged['base_price_Internal_mem']

for i in range(1):
    df_merged['Price'] += df_merged['base_price_RAM']

for i in range(1):
    df_merged['Price'] += df_merged['base_price_Gen_5G']

for i in range(1):
    df_merged['Price'] += df_merged['base_price_battery']

# df_merged = df_merged.drop(['Base_Price'], axis=1)
columns_to_drop = ['Base_Price', 'base_price_cpucore', 'base_price_dual_sim',
    ↳'base_price_Internal_mem', 'base_price_RAM', 'base_price_Gen_5G',
    ↳'base_price_battery']
df_merged.drop(columns=columns_to_drop, inplace=True)

```

[277]: df_merged

[277]:

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	Internal_mem(GB)	\
0	Samsung	180.32	312	1	1.5	No	16	
1	Huawei	159.12	362	1	1.6	No	64	
2	Google	159.29	241	1	1.0	No	16	
3	OnePlus	214.38	555	4	1.2	Yes	32	
4	Vivo	102.43	607	8	1.8	Yes	128	
...	
1495	LG	215.99	627	3	1.6	No	16	

1496	Apple	154.76	461	2	1.3	Yes	32
1497	Sony	211.70	466	4	1.4	No	256
1498	LG	100.70	742	2	1.7	Yes	32
1499	Xiaomi	219.00	677	4	1.5	No	256

	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	Price
0	8	77	31	Yes	5000	17	60500.0
1	8	51	13	No	3000	12	53400.0
2	11	84	23	No	5000	11	69920.0
3	7	91	10	No	6000	8	67070.0
4	12	71	27	No	4500	11	41130.0
...
1495	11	54	11	Yes	4500	15	32500.0
1496	4	63	25	No	5000	6	101010.0
1497	7	77	25	No	6000	13	43100.0
1498	6	43	31	No	3000	5	27920.0
1499	5	43	22	Yes	3500	6	34550.0

[1500 rows x 14 columns]

```
[278]: df_merged['Price'] = np.abs(df_merged['Price'])
df_merged['Price'] = df_merged['Price'].round(-3)
```

6 Describe about dataframe(df)

```
[279]: df_merged.describe()
```

```
[279]:
```

	Weight(gm)	PPI	CPU_core	CPU_freq	Internal_mem(GB)	\
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.000000	
mean	176.520353	449.583333	4.427333	1.451667	98.570667	
std	43.239296	200.174228	2.345006	0.287249	86.612660	
min	100.120000	100.000000	1.000000	1.000000	16.000000	
25%	139.862500	272.750000	2.000000	1.200000	32.000000	
50%	175.975000	451.000000	4.000000	1.500000	64.000000	
75%	214.347500	624.000000	6.000000	1.700000	128.000000	
max	249.940000	799.000000	8.000000	1.900000	256.000000	

	RAM	RearCam	Front_Cam	Battery	Thickness	\
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.000000	
mean	10.054667	63.491333	21.166667	4525.333333	11.478000	
std	3.786983	25.918573	6.680767	1006.992674	4.028938	
min	4.000000	20.000000	10.000000	3000.000000	5.000000	
25%	7.000000	41.000000	15.000000	3500.000000	8.000000	
50%	10.000000	63.000000	21.000000	4500.000000	11.000000	
75%	14.000000	86.000000	27.000000	5500.000000	15.000000	
max	16.000000	108.000000	32.000000	6000.000000	18.000000	

	Price
count	1500.000000
mean	52260.000000
std	22552.564054
min	22000.000000
25%	35000.000000
50%	42000.000000
75%	67000.000000
max	119000.000000

```
[280]: # df_merged['Price'] = np.where((df_merged['Dual_sim'] == 0) &
      ↪ (df_merged['Gen_5G'] == 0), df_merged['Price'] - 10000, df_merged['Price'])
      # df_merged['Price'] = np.where((df_merged['Dual_sim'] == 0) &
      ↪ (df_merged['Gen_5G'] == 1), df_merged['Price'] - 5000, df_merged['Price'])
      # df_merged['Price'] = np.where((df_merged['Dual_sim'] == 1) &
      ↪ (df_merged['Gen_5G'] == 1), df_merged['Price'], df_merged['Price'])
      # df_merged['Price'] = np.where((df_merged['Dual_sim'] == 1) &
      ↪ (df_merged['Gen_5G'] == 0), df_merged['Price'] - 5000, df_merged['Price'])
```

```
[281]: df_merged
```

```
[281]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	Internal_mem(GB)	\
0	Samsung	180.32	312	1	1.5	No	16	
1	Huawei	159.12	362	1	1.6	No	64	
2	Google	159.29	241	1	1.0	No	16	
3	OnePlus	214.38	555	4	1.2	Yes	32	
4	Vivo	102.43	607	8	1.8	Yes	128	
...	
1495	LG	215.99	627	3	1.6	No	16	
1496	Apple	154.76	461	2	1.3	Yes	32	
1497	Sony	211.70	466	4	1.4	No	256	
1498	LG	100.70	742	2	1.7	Yes	32	
1499	Xiaomi	219.00	677	4	1.5	No	256	

	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	Price
0	8	77	31	Yes	5000	17	60000.0
1	8	51	13	No	3000	12	53000.0
2	11	84	23	No	5000	11	70000.0
3	7	91	10	No	6000	8	67000.0
4	12	71	27	No	4500	11	41000.0
...	
1495	11	54	11	Yes	4500	15	32000.0
1496	4	63	25	No	5000	6	101000.0
1497	7	77	25	No	6000	13	43000.0
1498	6	43	31	No	3000	5	28000.0
1499	5	43	22	Yes	3500	6	35000.0

[1500 rows x 14 columns]

```
[282]: numeric_columns = df_merged.select_dtypes(include=[np.float, np.int]).columns.  
        ↳tolist()  
        non_numeric_columns = df_merged.select_dtypes(exclude=[np.number]).columns.  
        ↳tolist()  
        print("Numeric Columns:", numeric_columns)  
        print("Non-Numeric Columns:", non_numeric_columns)
```

Numeric Columns: ['Weight(gm)', 'PPI', 'CPU_core', 'CPU_freq',
'Internal_mem(GB)', 'RAM', 'RearCam', 'Front_Cam', 'Battery', 'Thickness',
'Price']
Non-Numeric Columns: ['Company', 'Dual_sim', 'Gen_5G']

```
[283]: csv_file_path = 'mobile_price_data.csv'  
        df_merged.to_csv(csv_file_path, index=False)  
        print(f"DataFrame has been successfully converted to CSV: {csv_file_path}")
```

DataFrame has been successfully converted to CSV: mobile_price_data.csv

7 finding Max and Min price for each company

```
[284]: # Assuming df is your DataFrame  
        company_list = df_merged['Company'].unique()  
  
        for company in company_list:  
            max_price = df_merged[df_merged['Company'] == company]['Price'].max()  
            min_price = df_merged[df_merged['Company'] == company]['Price'].min()  
            print(f"Company: {company}, Max Price: {max_price}, Min Price:   
        ↳{min_price}\n")
```

Company: Samsung, Max Price: 76000.0, Min Price: 56000.0

Company: Huawei, Max Price: 73000.0, Min Price: 51000.0

Company: Google, Max Price: 87000.0, Min Price: 67000.0

Company: OnePlus, Max Price: 78000.0, Min Price: 58000.0

Company: Vivo, Max Price: 49000.0, Min Price: 27000.0

Company: Motorola, Max Price: 40000.0, Min Price: 22000.0

Company: Apple, Max Price: 119000.0, Min Price: 95000.0

Company: Sony, Max Price: 53000.0, Min Price: 34000.0

Company: Xiaomi, Max Price: 46000.0, Min Price: 27000.0

Company: Nokia, Max Price: 47000.0, Min Price: 25000.0

Company: Realme, Max Price: 52000.0, Min Price: 29000.0

Company: LG, Max Price: 45000.0, Min Price: 26000.0

Company: Oppo, Max Price: 42000.0, Min Price: 23000.0

8 Performing label encoding (converting catogirical data to numerical data)

```
[285]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

# Define the columns to encode
columns_to_encode = ['Company', 'Gen_5G', 'Dual_sim']

# Create a dictionary to store encoding mappings
encoding_mappings = {}

# Encode each column and store the mappings
for column in columns_to_encode:
    df_merged[column] = label_encoder.fit_transform(df_merged[column]).
    ↪astype(int)

    # Store encoding mappings in the dictionary
    encoding_mapping = dict(zip(label_encoder.classes_, label_encoder.
    ↪transform(label_encoder.classes_)))
    encoding_mappings[column] = encoding_mapping

# Display the encoding mappings
for column, mappings in encoding_mappings.items():
    print(f"\nEncoding for {column}:")
    for category, encoded_number in mappings.items():
        print(f"{category}: {encoded_number}")
```

Encoding for Company:

Apple: 0

Google: 1

Huawei: 2

LG: 3
 Motorola: 4
 Nokia: 5
 OnePlus: 6
 Oppo: 7
 Realme: 8
 Samsung: 9
 Sony: 10
 Vivo: 11
 Xiaomi: 12

Encoding for Gen_5G:

No: 0

Yes: 1

Encoding for Dual_sim:

No: 0

Yes: 1

[286]: df_merged

```
[286]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	9	180.32	312	1	1.5	0	
1	2	159.12	362	1	1.6	0	
2	1	159.29	241	1	1.0	0	
3	6	214.38	555	4	1.2	1	
4	11	102.43	607	8	1.8	1	
...	
1495	3	215.99	627	3	1.6	0	
1496	0	154.76	461	2	1.3	1	
1497	10	211.70	466	4	1.4	0	
1498	3	100.70	742	2	1.7	1	
1499	12	219.00	677	4	1.5	0	

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	\
0	16	8	77	31	1	5000	17	
1	64	8	51	13	0	3000	12	
2	16	11	84	23	0	5000	11	
3	32	7	91	10	0	6000	8	
4	128	12	71	27	0	4500	11	
...	
1495	16	11	54	11	1	4500	15	
1496	32	4	63	25	0	5000	6	
1497	256	7	77	25	0	6000	13	
1498	32	6	43	31	0	3000	5	
1499	256	5	43	22	1	3500	6	

	Price
0	60000.0
1	53000.0
2	70000.0
3	67000.0
4	41000.0
...	...
1495	32000.0
1496	101000.0
1497	43000.0
1498	28000.0
1499	35000.0

[1500 rows x 14 columns]

```
[287]: df_merged['Weight(gm)'] = df_merged['Weight(gm)'].round(-1)
df_merged.head()
```

```
[287]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	Internal_mem(GB)	\
0	9	180.0	312	1	1.5	0	16	
1	2	160.0	362	1	1.6	0	64	
2	1	160.0	241	1	1.0	0	16	
3	6	210.0	555	4	1.2	1	32	
4	11	100.0	607	8	1.8	1	128	


	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	Price
0	8	77	31	1	5000	17	60000.0
1	8	51	13	0	3000	12	53000.0
2	11	84	23	0	5000	11	70000.0
3	7	91	10	0	6000	8	67000.0
4	12	71	27	0	4500	11	41000.0

9 Inserting some null values

```
[288]: import numpy as np
df_merged.loc[df.sample(130).index, 'PPI'] = np.nan
df_merged.loc[df.sample(90).index, 'Internal_mem(GB)'] = np.nan
df_merged.loc[df.sample(50).index, 'Battery'] = np.nan
df_merged.loc[df.sample(30).index, 'CPU_core'] = np.nan
df_merged.loc[df.sample(30).index, 'Gen_5G'] = np.nan
```

```
[289]: import missingno as msno
msno.matrix(df_merged)
```

```
[289]: <Axes: >
```



```
[290]: df_merged.isnull().sum()
```

```
[290]: Company          0
Weight(gm)            0
PPI                  130
CPU_core              30
CPU_freq              0
Dual_sim              0
Internal_mem(GB)      90
RAM                   0
RearCam               0
Front_Cam             0
Gen_5G                30
Battery              50
Thickness             0
Price                 0
dtype: int64
```

```
[291]: df_merged
```

```
[291]:   Company  Weight(gm)  PPI  CPU_core  CPU_freq  Dual_sim  \
0         9      180.0  312.0         1.0        1.5         0
1         2      160.0  362.0         1.0        1.6         0
2         1      160.0  241.0         1.0        1.0         0
3         6      210.0  555.0         4.0        1.2         1
4        11      100.0  607.0         8.0        1.8         1
...     ...      ...    ...      ...      ...      ...
1495      3      220.0  627.0         3.0        1.6         0
1496      0      150.0  461.0         2.0        1.3         1
```

1497	10	210.0	466.0	4.0	1.4	0
1498	3	100.0	742.0	2.0	1.7	1
1499	12	220.0	677.0	4.0	1.5	0

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	\
0	16.0	8	77	31	1.0	5000.0	17	
1	64.0	8	51	13	0.0	3000.0	12	
2	16.0	11	84	23	0.0	5000.0	11	
3	32.0	7	91	10	0.0	6000.0	8	
4	128.0	12	71	27	0.0	4500.0	11	
...	
1495	16.0	11	54	11	1.0	4500.0	15	
1496	32.0	4	63	25	NaN	5000.0	6	
1497	NaN	7	77	25	0.0	6000.0	13	
1498	NaN	6	43	31	0.0	3000.0	5	
1499	256.0	5	43	22	1.0	3500.0	6	

	Price
0	60000.0
1	53000.0
2	70000.0
3	67000.0
4	41000.0
...	...
1495	32000.0
1496	101000.0
1497	43000.0
1498	28000.0
1499	35000.0

[1500 rows x 14 columns]

```
[292]: df_merged['Company'].value_counts()
```

```
[292]: 2    144
0     128
4     120
9     116
10    116
11    115
12    113
1     110
8     109
6     108
3     108
5     107
7     106
```


Name: Company, dtype: int64

10 Imputation of missing values using various techniques

```
[293]: # from sklearn.experimental import enable_iterative_imputer
# from sklearn.impute import IterativeImputer

# # Assuming df is your DataFrame and 'Company' is the column you want to impute
# imputer_multiple = IterativeImputer(random_state=42)
# df_merged['Company'] = imputer_multiple.fit_transform(df_merged[['Company']])
```

```
[294]: # from sklearn.impute import KNNImputer
# imputer_knn = KNNImputer(n_neighbors=5)
# df_merged['Company'] = imputer_knn.fit_transform(df_merged[['Company']]).
#     ↳ astype(int)
# df_merged
```

11 Using median

```
[295]: import pandas as pd
from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy='median')
df_merged['Battery'] = imputer.fit_transform(df_merged[['Battery']])

df_merged
```

```
[295]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	9	180.0	312.0	1.0	1.5	0	
1	2	160.0	362.0	1.0	1.6	0	
2	1	160.0	241.0	1.0	1.0	0	
3	6	210.0	555.0	4.0	1.2	1	
4	11	100.0	607.0	8.0	1.8	1	
...		
1495	3	220.0	627.0	3.0	1.6	0	
1496	0	150.0	461.0	2.0	1.3	1	
1497	10	210.0	466.0	4.0	1.4	0	
1498	3	100.0	742.0	2.0	1.7	1	
1499	12	220.0	677.0	4.0	1.5	0	

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	\
0	16.0	8	77	31	1.0	5000.0	17	
1	64.0	8	51	13	0.0	3000.0	12	
2	16.0	11	84	23	0.0	5000.0	11	
3	32.0	7	91	10	0.0	6000.0	8	

4	128.0	12	71	27	0.0	4500.0	11
...
1495	16.0	11	54	11	1.0	4500.0	15
1496	32.0	4	63	25	NaN	5000.0	6
1497	NaN	7	77	25	0.0	6000.0	13
1498	NaN	6	43	31	0.0	3000.0	5
1499	256.0	5	43	22	1.0	3500.0	6

	Price
0	60000.0
1	53000.0
2	70000.0
3	67000.0
4	41000.0
...	...
1495	32000.0
1496	101000.0
1497	43000.0
1498	28000.0
1499	35000.0

[1500 rows x 14 columns]

12 Using Mode

```
[296]: imputer = SimpleImputer(strategy='most_frequent')
df_merged['Internal_mem(GB)'] = imputer.
        ↪fit_transform(df_merged[['Internal_mem(GB)']])
df_merged
```

```
[296]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	9	180.0	312.0	1.0	1.5	0	
1	2	160.0	362.0	1.0	1.6	0	
2	1	160.0	241.0	1.0	1.0	0	
3	6	210.0	555.0	4.0	1.2	1	
4	11	100.0	607.0	8.0	1.8	1	
...
1495	3	220.0	627.0	3.0	1.6	0	
1496	0	150.0	461.0	2.0	1.3	1	
1497	10	210.0	466.0	4.0	1.4	0	
1498	3	100.0	742.0	2.0	1.7	1	
1499	12	220.0	677.0	4.0	1.5	0	

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	\
0	16.0	8	77	31	1.0	5000.0	17	
1	64.0	8	51	13	0.0	3000.0	12	

2	16.0	11	84	23	0.0	5000.0	11
3	32.0	7	91	10	0.0	6000.0	8
4	128.0	12	71	27	0.0	4500.0	11
...
1495	16.0	11	54	11	1.0	4500.0	15
1496	32.0	4	63	25	NaN	5000.0	6
1497	16.0	7	77	25	0.0	6000.0	13
1498	16.0	6	43	31	0.0	3000.0	5
1499	256.0	5	43	22	1.0	3500.0	6

	Price
0	60000.0
1	53000.0
2	70000.0
3	67000.0
4	41000.0
...	...
1495	32000.0
1496	101000.0
1497	43000.0
1498	28000.0
1499	35000.0

[1500 rows x 14 columns]

13 Using Knn

```
[297]: from sklearn.impute import KNNImputer
imputer_knn = KNNImputer(n_neighbors=5)
df_merged['PPI'] = imputer_knn.fit_transform(df_merged[['PPI']]).astype(int)
df_merged
```

```
[297]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	9	180.0	312	1.0	1.5	0	
1	2	160.0	362	1.0	1.6	0	
2	1	160.0	241	1.0	1.0	0	
3	6	210.0	555	4.0	1.2	1	
4	11	100.0	607	8.0	1.8	1	
...
1495	3	220.0	627	3.0	1.6	0	
1496	0	150.0	461	2.0	1.3	1	
1497	10	210.0	466	4.0	1.4	0	
1498	3	100.0	742	2.0	1.7	1	
1499	12	220.0	677	4.0	1.5	0	

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	\
--	------------------	-----	---------	-----------	--------	---------	-----------	---

0	16.0	8	77	31	1.0	5000.0	17
1	64.0	8	51	13	0.0	3000.0	12
2	16.0	11	84	23	0.0	5000.0	11
3	32.0	7	91	10	0.0	6000.0	8
4	128.0	12	71	27	0.0	4500.0	11
...
1495	16.0	11	54	11	1.0	4500.0	15
1496	32.0	4	63	25	NaN	5000.0	6
1497	16.0	7	77	25	0.0	6000.0	13
1498	16.0	6	43	31	0.0	3000.0	5
1499	256.0	5	43	22	1.0	3500.0	6

	Price
0	60000.0
1	53000.0
2	70000.0
3	67000.0
4	41000.0
...	...
1495	32000.0
1496	101000.0
1497	43000.0
1498	28000.0
1499	35000.0

[1500 rows x 14 columns]

14 Multivariate feature imputation

```
[298]: from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
mice_imputer = IterativeImputer(random_state=0)
df_merged['CPU_core'] = mice_imputer.fit_transform(df_merged[['CPU_core']]).
    .astype(int)
df_merged
```

```
[298]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	9	180.0	312	1	1.5	0	
1	2	160.0	362	1	1.6	0	
2	1	160.0	241	1	1.0	0	
3	6	210.0	555	4	1.2	1	
4	11	100.0	607	8	1.8	1	
...
1495	3	220.0	627	3	1.6	0	
1496	0	150.0	461	2	1.3	1	
1497	10	210.0	466	4	1.4	0	

1498	3	100.0	742	2	1.7	1
1499	12	220.0	677	4	1.5	0

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness \
0	16.0	8	77	31	1.0	5000.0	17
1	64.0	8	51	13	0.0	3000.0	12
2	16.0	11	84	23	0.0	5000.0	11
3	32.0	7	91	10	0.0	6000.0	8
4	128.0	12	71	27	0.0	4500.0	11
...
1495	16.0	11	54	11	1.0	4500.0	15
1496	32.0	4	63	25	NaN	5000.0	6
1497	16.0	7	77	25	0.0	6000.0	13
1498	16.0	6	43	31	0.0	3000.0	5
1499	256.0	5	43	22	1.0	3500.0	6

	Price
0	60000.0
1	53000.0
2	70000.0
3	67000.0
4	41000.0
...	...
1495	32000.0
1496	101000.0
1497	43000.0
1498	28000.0
1499	35000.0

[1500 rows x 14 columns]

```
[299]: df_merged['CPU_core'].value_counts()
```

```
[299]: 1    219
        6    199
        4    196
        7    190
        2    180
        3    178
        8    175
        5    163
        Name: CPU_core, dtype: int64
```


```
[300]: df_merged['Gen_5G'] = df_merged['Gen_5G'].fillna(method='bfill')
```

```
[301]: df_merged.isnull().sum()
```

```
[301]: Company          0
      Weight(gm)        0
      PPI                0
      CPU_core           0
      CPU_freq           0
      Dual_sim           0
      Internal_mem(GB)    0
      RAM                0
      RearCam            0
      Front_Cam          0
      Gen_5G             0
      Battery            0
      Thickness          0
      Price              0
      dtype: int64
```

```
[302]: msno.matrix(df_merged)
```

```
[302]: <Axes: >
```



```
[303]: df_merged.corr()
```

```
[303]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	\
Company	1.000000	-0.033894	-0.035367	-0.001696	-0.043763	
Weight(gm)	-0.033894	1.000000	-0.024519	0.044470	-0.004677	
PPI	-0.035367	-0.024519	1.000000	0.024396	0.055434	
CPU_core	-0.001696	0.044470	0.024396	1.000000	-0.008669	
CPU_freq	-0.043763	-0.004677	0.055434	-0.008669	1.000000	
Dual_sim	-0.016431	0.008557	0.019944	-0.003378	-0.007419	

Internal_mem(GB)	0.007243	0.017536	0.003930	0.002807	0.037610
RAM	0.024401	-0.025259	0.019059	0.005788	0.047505
RearCam	-0.026918	0.027723	0.025406	0.009979	-0.038035
Front_Cam	0.013403	0.042453	0.042920	0.023586	-0.002926
Gen_5G	-0.014889	0.021557	0.031871	0.018104	-0.014477
Battery	-0.001666	0.023796	-0.008695	0.016072	0.005073
Thickness	0.037313	0.011519	0.018738	0.019412	-0.035592
Price	-0.556091	0.034972	0.056383	0.039368	0.005700

	Dual_sim	Internal_mem(GB)	RAM	RearCam	Front_Cam	\
Company	-0.016431	0.007243	0.024401	-0.026918	0.013403	
Weight(gm)	0.008557	0.017536	-0.025259	0.027723	0.042453	
PPI	0.019944	0.003930	0.019059	0.025406	0.042920	
CPU_core	-0.003378	0.002807	0.005788	0.009979	0.023586	
CPU_freq	-0.007419	0.037610	0.047505	-0.038035	-0.002926	
Dual_sim	1.000000	-0.023921	-0.010760	0.027898	0.007459	
Internal_mem(GB)	-0.023921	1.000000	0.054864	-0.021540	-0.006787	
RAM	-0.010760	0.054864	1.000000	-0.034495	-0.033373	
RearCam	0.027898	-0.021540	-0.034495	1.000000	-0.043369	
Front_Cam	0.007459	-0.006787	-0.033373	-0.043369	1.000000	
Gen_5G	-0.007012	0.029458	0.014478	0.015769	0.032820	
Battery	-0.027663	-0.037936	0.047197	0.003270	-0.033193	
Thickness	-0.035409	0.016175	0.046950	0.026683	0.024301	
Price	0.050466	0.054010	0.056190	0.042059	0.006234	

	Gen_5G	Battery	Thickness	Price
Company	-0.014889	-0.001666	0.037313	-0.556091
Weight(gm)	0.021557	0.023796	0.011519	0.034972
PPI	0.031871	-0.008695	0.018738	0.056383
CPU_core	0.018104	0.016072	0.019412	0.039368
CPU_freq	-0.014477	0.005073	-0.035592	0.005700
Dual_sim	-0.007012	-0.027663	-0.035409	0.050466
Internal_mem(GB)	0.029458	-0.037936	0.016175	0.054010
RAM	0.014478	0.047197	0.046950	0.056190
RearCam	0.015769	0.003270	0.026683	0.042059
Front_Cam	0.032820	-0.033193	0.024301	0.006234
Gen_5G	1.000000	0.052340	0.009293	0.078342
Battery	0.052340	1.000000	-0.037837	0.048829
Thickness	0.009293	-0.037837	1.000000	-0.001714
Price	0.078342	0.048829	-0.001714	1.000000

```
[304]: df_merged['Internal_mem(GB)'].value_counts()
```

```
[304]: 16.0      387
        64.0      290
        128.0     289
        256.0     283
```

```
32.0      251
Name: Internal_mem(GB), dtype: int64
```

```
[305]: df_merged['Battery'].value_counts()
```

```
[305]: 5000.0      234
      4500.0      224
      5500.0      215
      6000.0      213
      4000.0      208
      3000.0      204
      3500.0      202
      Name: Battery, dtype: int64
```

15 Data Quality


```
[306]: import pandas as pd
      duplicate_rows = df_merged[df_merged.duplicated()]
      duplicate_rows.shape
```

```
[306]: (0, 14)
```

So there are no duplicates

16 HeatMap

```
[307]: plt.figure(figsize=(20, 8))
      correlation_matrix = df_merged.corr()
      sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
      plt.show()
```

17 Data Skewness


```
[308]: numeric_columns = df_merged.select_dtypes(include=['float64', 'int']).columns
skewness = df_merged[numeric_columns].apply(lambda x: x.skew())
print(skewness)
```

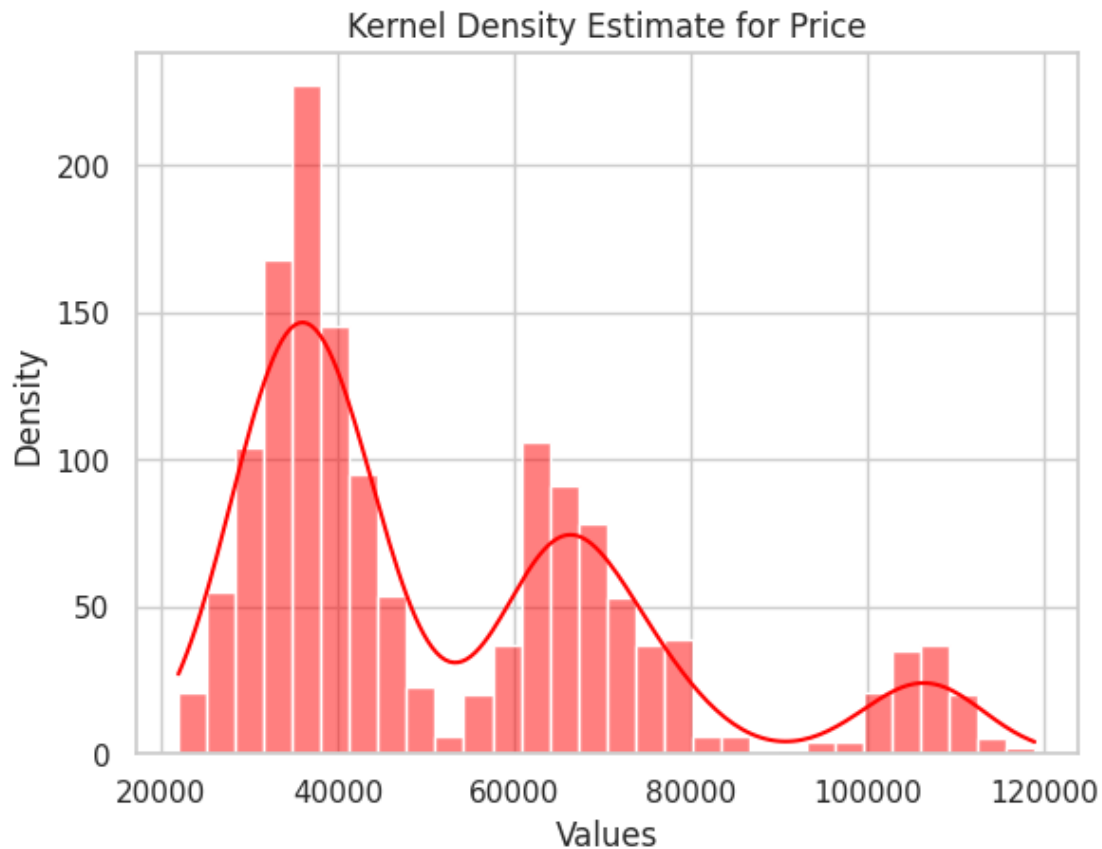
```
Company          0.040592
Weight(gm)       -0.052652
PPI              0.015320
CPU_core         0.009454
CPU_freq        -0.017846
Dual_sim        -0.074794
Internal_mem(GB)  0.949655
RAM             -0.025163
RearCam          0.025589
Front_Cam       -0.022150
Gen_5G           0.053406
Battery         -0.047799
Thickness        0.019966
Price            1.057614
dtype: float64
```

```
[309]: columns_to_analyze = ["Price"]
for column in columns_to_analyze:
    sns.histplot(df_merged[column], kde=True, color='blue', bins=30)
    plt.title(f'Distribution of {column}')
    plt.xlabel('Values')
```

```
plt.ylabel('Frequency')
plt.show()

sns.histplot(df_merged[column], kde=True, color='red', bins=30)
plt.title(f'Kernel Density Estimate for {column}')
plt.xlabel('Values')
plt.ylabel('Density')
plt.show()
```





```
[310]: dfcopy = df_merged.copy()
```

18 Skneness after Transformation

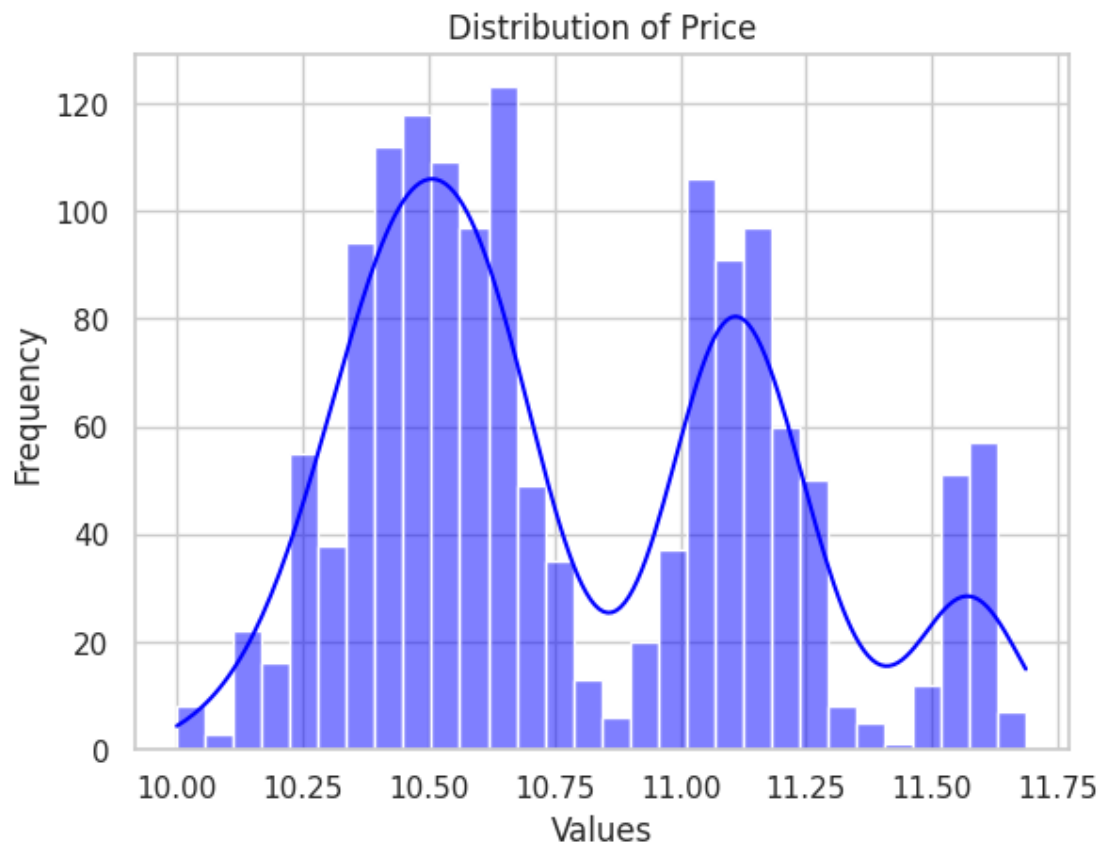
```
[311]: columns_to_analyze = ["Price"]
for column in columns_to_analyze:
    df_merged[column] = df_merged[column].apply(lambda x: np.log1p(x) if x > 0
    ↪ else 0)
    transformed_skewness = df_merged[column].skew()
    print(f"Skewness after transformation for {column}: {transformed_skewness}")
```

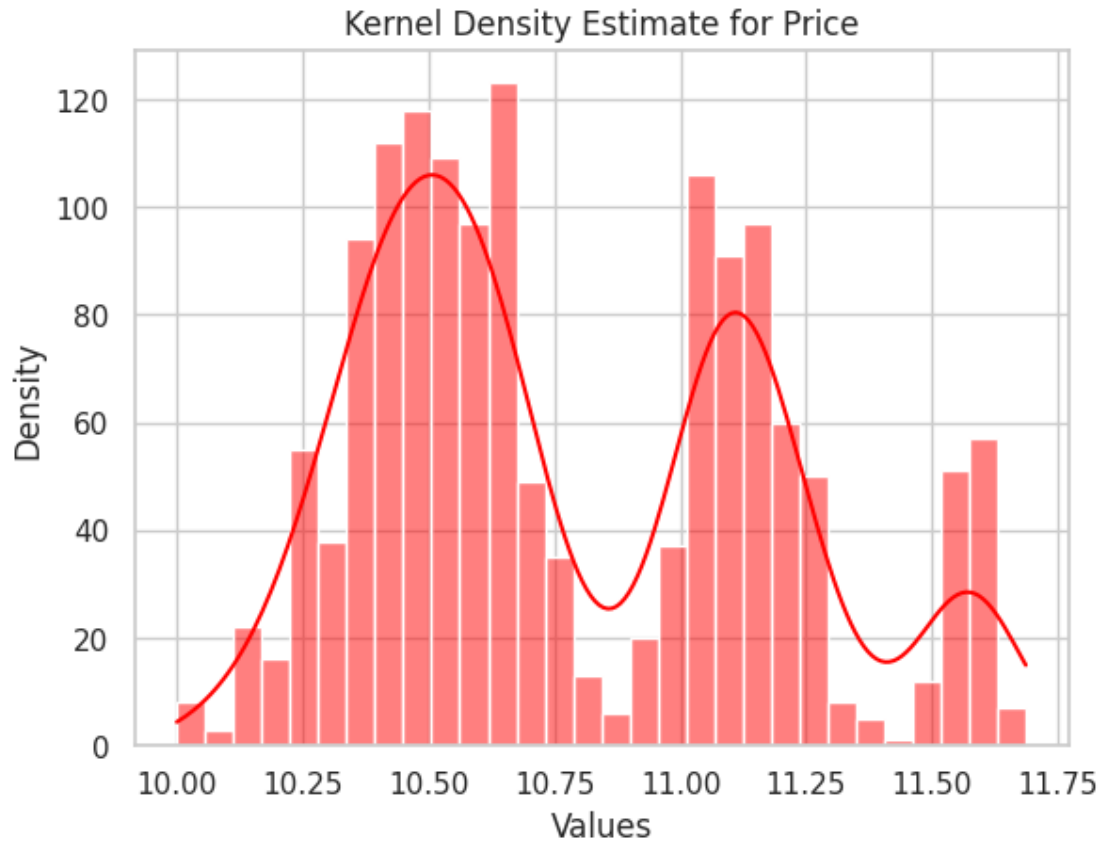
Skewness after transformation for Price: 0.4522692673157849

```
[312]: columns_to_analyze = ["Price"]
for column in columns_to_analyze:
    sns.histplot(df_merged[column], kde=True, color='blue', bins=30)
    plt.title(f'Distribution of {column}')
    plt.xlabel('Values')
    plt.ylabel('Frequency')
```

```
plt.show()

sns.histplot(df_merged[column], kde=True, color='red', bins=30)
plt.title(f'Kernel Density Estimate for {column}')
plt.xlabel('Values')
plt.ylabel('Density')
plt.show()
```





19 Normalization

```
[313]: df_Normalize=df_merged.copy()
```

```
[314]: from sklearn.preprocessing import MinMaxScaler
columns_to_normalize = ['Price']
scaler = MinMaxScaler()
df_Normalize[columns_to_normalize] = scaler.
    ↪fit_transform(df_Normalize[columns_to_normalize])
df_Normalize.head(n=3)
```

```
[314]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	Internal_mem(GB)	\
0	9	180.0	312	1	1.5	0	16.0	
1	2	160.0	362	1	1.6	0	64.0	
2	1	160.0	241	1	1.0	0	16.0	

	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	Price
0	8	77	31	1.0	5000.0	17	0.594341
1	8	51	13	0.0	3000.0	12	0.520853

2	11	84	23	0.0	5000.0	11	0.685658
---	----	----	----	-----	--------	----	----------

```
[315]: df_merged=df_Normalize.copy()
```

20 Binning

```
[316]: import pandas as pd
column_name = 'Price'
num_bins = 10

df_Normalize['Binned_Column'] = pd.cut(df_Normalize[column_name],
    ↪bins=num_bins, labels=False)

binned_means = df_Normalize.groupby('Binned_Column')[column_name].mean()
binned_medians = df_Normalize.groupby('Binned_Column')[column_name].median()

df_Normalize['Binned_Column_Means'] = df_Normalize['Binned_Column'].
    ↪map(binned_means)
df_Normalize['Binned_Column_Medians'] = df_Normalize['Binned_Column'].
    ↪map(binned_medians)

df_Normalize[['Binned_Column', 'Binned_Column_Means', 'Binned_Column_Medians',
    ↪column_name]]
```

```
[316]:
```

	Binned_Column	Binned_Column_Means	Binned_Column_Medians	Price
0	5	0.570393	0.574258	0.594341
1	5	0.570393	0.574258	0.520853
2	6	0.649945	0.650802	0.685658
3	6	0.649945	0.650802	0.659710
4	3	0.348768	0.354147	0.368775
...
1495	2	0.250888	0.257874	0.221961
1496	9	0.937023	0.937031	0.902845
1497	3	0.348768	0.354147	0.396989
1498	1	0.159285	0.163646	0.142859
1499	2	0.250888	0.257874	0.275045

[1500 rows x 4 columns]

21 Equi width

```
[317]: import pandas as pd
import matplotlib.pyplot as plt

column_name = 'Price'
```

```

fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))

# Original column values
axes[0].bar(df_Normalize.index, df_Normalize[column_name], color='green',
            alpha=0.5)
axes[0].set_title('Original Values')


# Binned means
axes[1].bar(df_Normalize['Binned_Column'].unique(), df_Normalize.
            groupby('Binned_Column')[column_name].mean(), color='red', alpha=0.5)
axes[1].set_title('Binned Means')

# Binned medians
axes[2].bar(df_Normalize['Binned_Column'].unique(), df_Normalize.
            groupby('Binned_Column')[column_name].median(), color='blue', alpha=0.5)
axes[2].set_title('Binned Medians')

# Set common labels
for ax in axes:
    ax.set_xlabel('Bins')
    ax.set_ylabel(column_name)

plt.tight_layout()
plt.show()

```



22 Equi depth

```

[318]: import pandas as pd
        column_name = 'Price'
        num_bins = 5
        # Calculate bin edges using pandas' qcut function

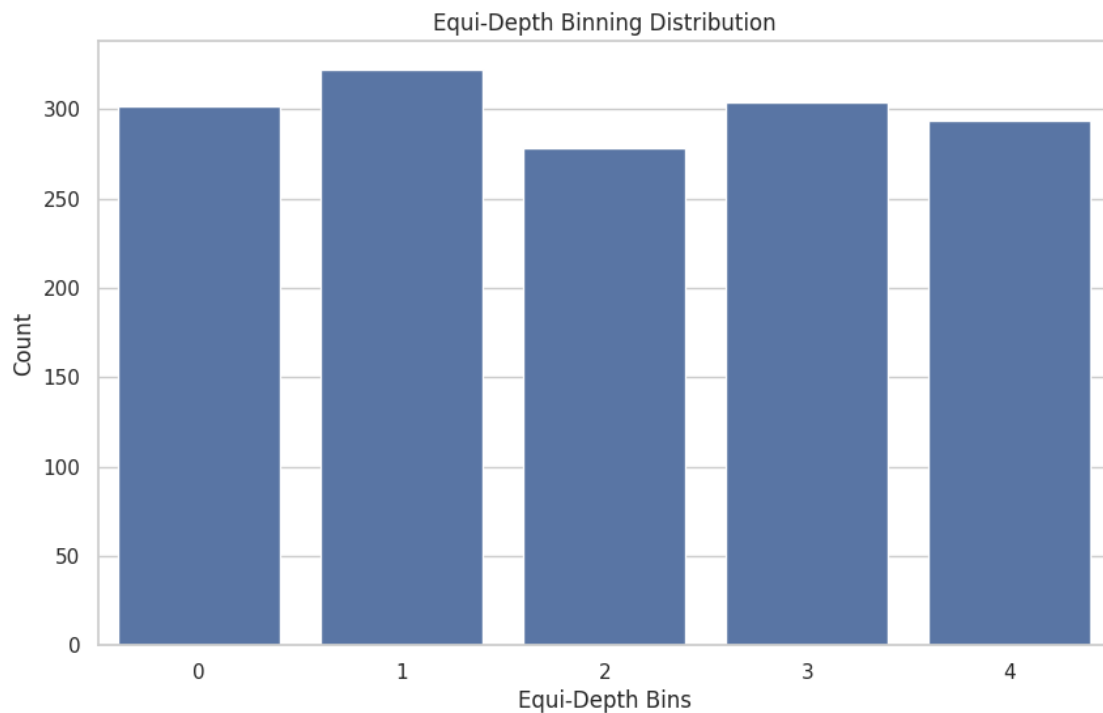
```

```
df_Normalize['EquiDepth_Binned_Column'], bin_edges = pd.
    qcut(df_Normalize[column_name], q=num_bins, retbins=True, labels=False)
print("Bin Edges:", bin_edges)
df_Normalize[['EquiDepth_Binned_Column', column_name]].head()
```

Bin Edges: [0. 0.24018921 0.33914956 0.54279596 0.67713466 1.]

```
[318]: EquiDepth_Binned_Column    Price
0                3  0.594341
1                2  0.520853
2                4  0.685658
3                3  0.659710
4                2  0.368775
```


```
[319]: sns.set(style="whitegrid")
# Create a bar plot to visualize the distribution of the equi-depth bins
plt.figure(figsize=(10, 6))
sns.countplot(x='EquiDepth_Binned_Column', data=df_Normalize)
plt.xlabel('Equi-Depth Bins')
plt.ylabel('Count')
plt.title('Equi-Depth Binning Distribution')
plt.show()
```






```
[320]: df_Normalize.drop('Binned_Column', axis=1, inplace=True)
df_Normalize.drop('Binned_Column_Means', axis=1, inplace=True)
df_Normalize.drop('Binned_Column_Medians', axis=1, inplace=True)
df_Normalize.drop('EquiDepth_Binned_Column', axis=1, inplace=True)
```


23 Outlier Detection


```
[321]: for column in df_merged.columns:
    plt.figure(figsize=(8, 6))
    sns.boxplot(x=column, data=df_Normalize)
    plt.xlabel(column)
    plt.ylabel('Count')
    plt.title('Box Plot')
    plt.show()
```







































```
[322]: plt.figure(figsize=(35, 8))
sns.boxplot(data=df_Normalize)
plt.title('Distribution of Laptop Prices')
plt.xticks(rotation='vertical')
plt.show()
```



```
[323]: sns.boxplot(data=df_Normalize, y='Price')
plt.title('Distribution of Prices')
plt.show()
```



```
[324]: max_weight = df_merged['Weight(gm)'].max()
min_weight = df_merged['Weight(gm)'].min()

print(f"Maximum Weight: {max_weight}")
print(f"Minimum Weight: {min_weight}")
```

```
Maximum Weight: 250.0
Minimum Weight: 100.0
```

24 Outlier detection for Price (target Column) using IQR Method

```
[325]: columns_with_outliers = ['Price']

def remove_outliers_iqr(df, columns):
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
```

```

IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

df = df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]

return df

IQRout_dataset = remove_outliers_iqr(df_Normalize, columns_with_outliers)

```

```

[326]: import pandas as pd
import matplotlib.pyplot as plt

p_column = df_Normalize['Price']
Q1 = p_column.quantile(0.25)
Q3 = p_column.quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
outliers = df_merged[(p_column < lower_bound) | (p_column > upper_bound)]
# Display the outliers
print("Outliers:")
print(outliers['Price'])


plt.boxplot(p_column)
plt.xlabel('Price')
plt.title('Outlier Detection using IQR')
plt.show()

```

```

Outliers:
Series([], Name: Price, dtype: float64)

```

25 Removing Outliers using Z_score method by Setting threshold Value

```
[327]: from scipy.stats import zscore
z_scores = zscore(df_merged)
abs_z_scores = np.abs(z_scores)
threshold = 2.5
df_no_outliers2 = df_Normalize[(abs_z_scores < threshold).all(axis=1)]
df_no_outliers2
```

```
[327]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	\
0	9	180.0	312	1	1.5	0	
1	2	160.0	362	1	1.6	0	
2	1	160.0	241	1	1.0	0	
3	6	210.0	555	4	1.2	1	
4	11	100.0	607	8	1.8	1	
...	
1495	3	220.0	627	3	1.6	0	

1496	0	150.0	461	2	1.3	1
1497	10	210.0	466	4	1.4	0
1498	3	100.0	742	2	1.7	1
1499	12	220.0	677	4	1.5	0

	Internal_mem(GB)	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	\
0	16.0	8	77	31	1.0	5000.0	17	
1	64.0	8	51	13	0.0	3000.0	12	
2	16.0	11	84	23	0.0	5000.0	11	
3	32.0	7	91	10	0.0	6000.0	8	
4	128.0	12	71	27	0.0	4500.0	11	
...	
1495	16.0	11	54	11	1.0	4500.0	15	
1496	32.0	4	63	25	0.0	5000.0	6	
1497	16.0	7	77	25	0.0	6000.0	13	
1498	16.0	6	43	31	0.0	3000.0	5	
1499	256.0	5	43	22	1.0	3500.0	6	

	Price
0	0.594341
1	0.520853
2	0.685658
3	0.659710
4	0.368775
...	...
1495	0.221961
1496	0.902845
1497	0.396989
1498	0.142859
1499	0.275045

[1500 rows x 14 columns]

```
[328]: df_no_outliers2.shape
```


```
[328]: (1500, 14)
```

```
[329]: df_no_outliers2['Company'].value_counts()
```

```
[329]: 2      144
0      128
4      120
9      116
10     116
11     115
12     113
1      110
```

```
8      109
6      108
3      108
5      107
7      106
Name: Company, dtype: int64
```


```
[330]: plt.figure(figsize=(15, 8))
sns.boxplot(data=df_no_outliers2)
plt.title('Outliers for each column')
plt.xticks(rotation='vertical')
plt.show()
```



```
[331]: sns.boxplot(data=df_no_outliers2, y='Price')
plt.title('Distribution of Laptop Prices')
plt.show()
```




```
[332]: plt.figure(figsize=(20, 8))
correlation_matrix = df_no_outliers2.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.show()
```

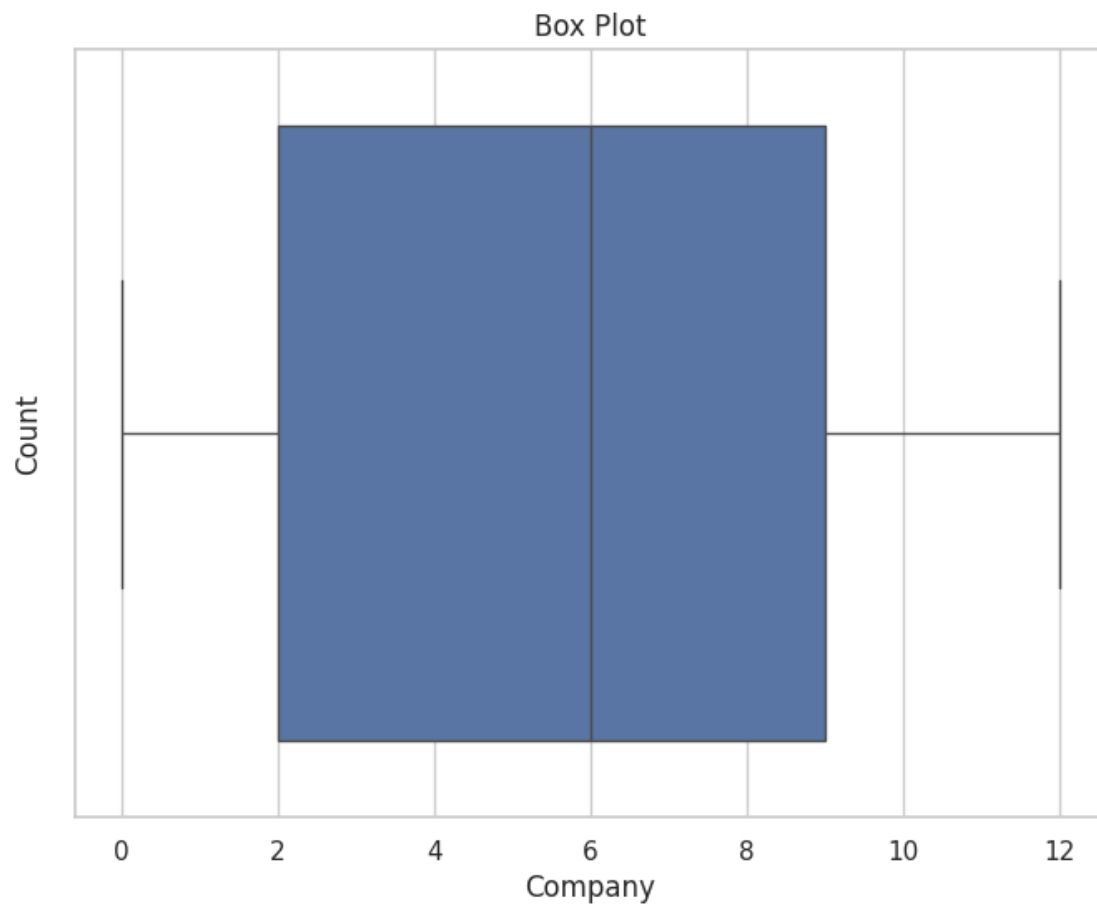



```
[333]: df_1 = df_merged.copy()
```

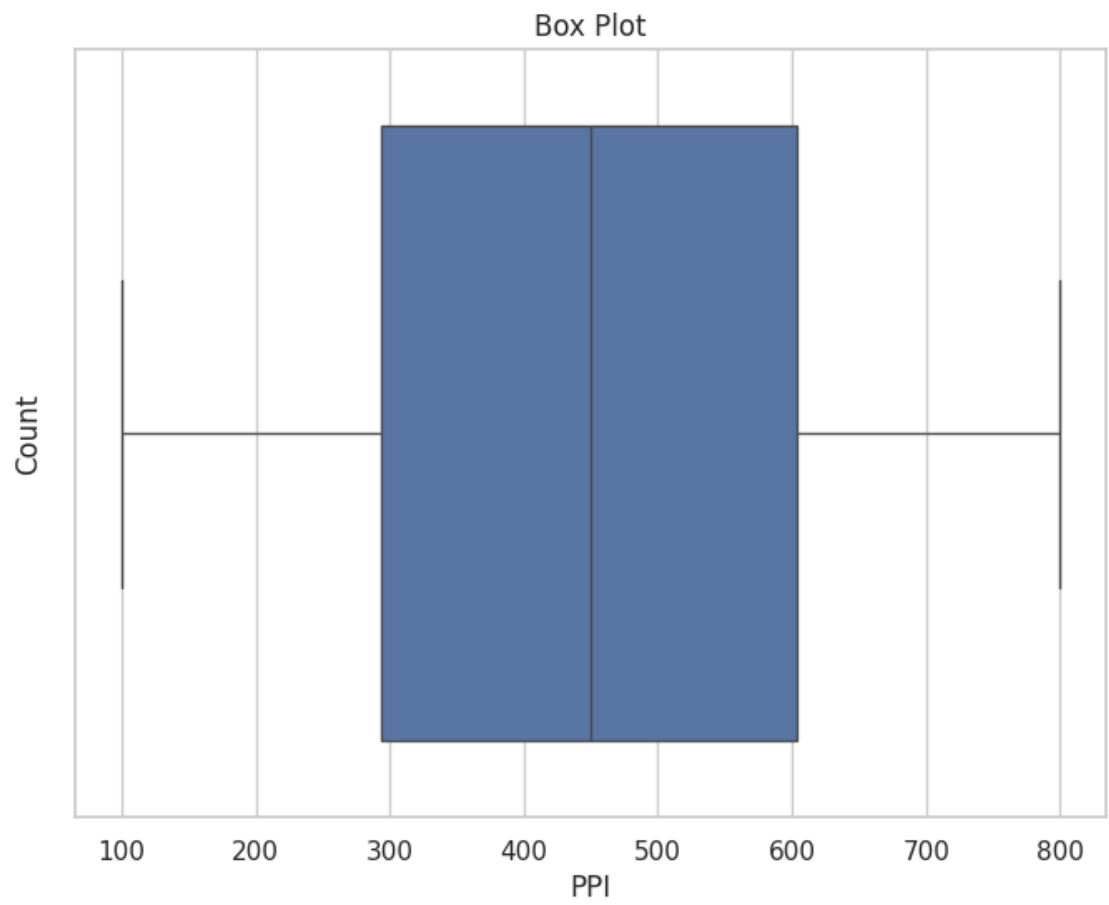
```
[334]: plt.figure(figsize=(12, 8))  
plt.xticks(rotation='vertical')  
sns.boxplot(data=df_no_outliers2)  
plt.title('Outliers for each column')  
plt.show()
```





```
[335]: for column in df_no_outliers2.columns:  
    plt.figure(figsize=(8, 6))  
    sns.boxplot(x=column, data=df_no_outliers2)  
    plt.xlabel(column)  
    plt.ylabel('Count')  
    plt.title('Box Plot')  
    plt.show()
```










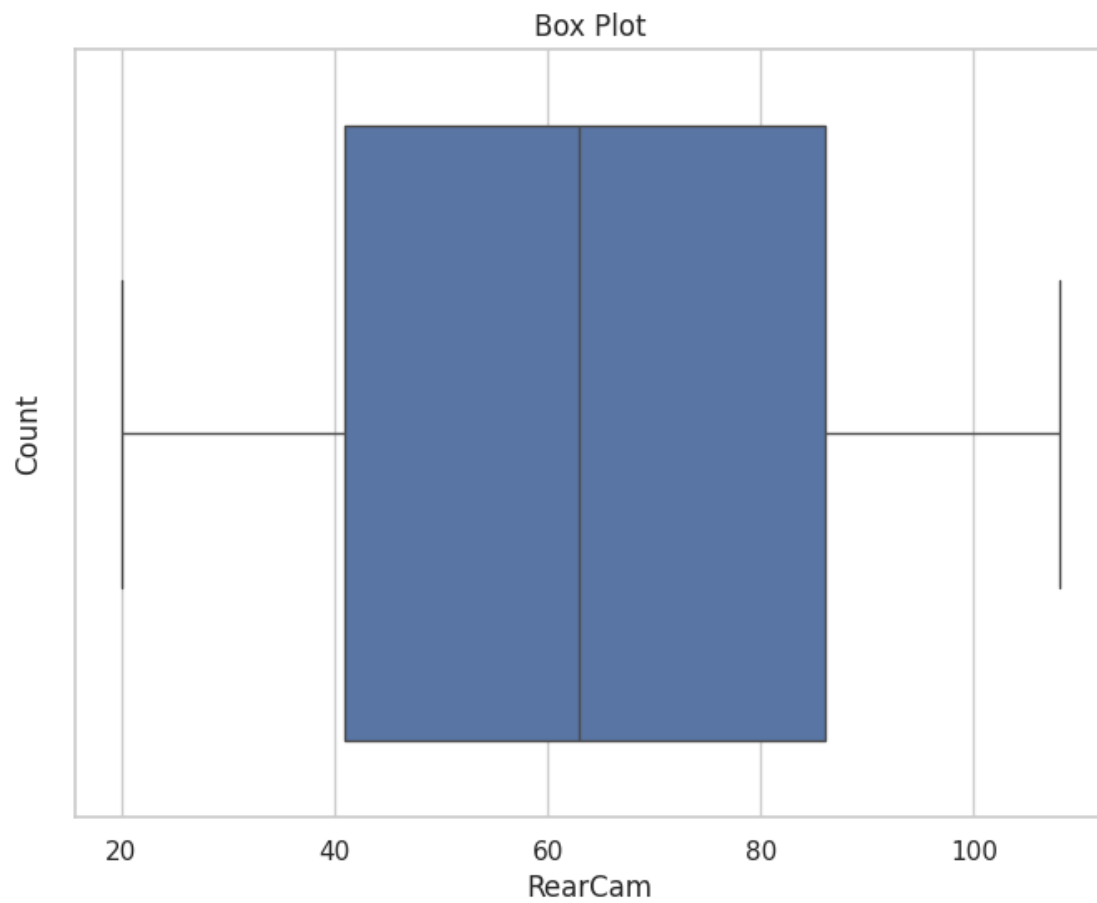



























```
[336]: df_no_outliers2.head(n=3)
```

```
[336]:   Company  Weight(gm)  PPI  CPU_core  CPU_freq  Dual_sim  Internal_mem(GB)  \
0        9      180.0   312         1      1.5         0         16.0
1        2      160.0   362         1      1.6         0         64.0
2        1      160.0   241         1      1.0         0         16.0
```

```
   RAM  RearCam  Front_Cam  Gen_5G  Battery  Thickness  Price
0    8       77        31     1.0   5000.0         17  0.594341
1    8       51        13     0.0   3000.0         12  0.520853
2   11       84        23     0.0   5000.0         11  0.685658
```

```
[372]: column_types = df_no_outliers2.dtypes
print(column_types)
```

```
Company          category
Weight(gm)       category
PPI              category
CPU_core         category
```

```

CPU_freq          category
Dual_sim          category
Internal_mem(GB)  category
RAM              category
RearCam           category
Front_Cam         category
Gen_5G           category
Battery           category
Thickness         category
Price             category
dtype: object

```


```


[337]: corr = df_no_outliers2.corr()
plt.figure(figsize=(12, 10))

sns.heatmap(corr[(corr <= 0.5) | (corr >= -0.4)], cmap='viridis', vmax=1.0, vmin=-1.0, linewidths=0.1,
            annot=True, annot_kws={"size": 8}, square=True);
corr = df_no_outliers2.drop('Price', axis=1).corr()
plt.figure(figsize=(12, 10))

sns.heatmap(corr[(corr <= 0.5) | (corr >= -0.4)],
            cmap='viridis', vmax=1.0, vmin=-1.0, linewidths=0.1,
            annot=True, annot_kws={"size": 8}, square=True);

```





26 EDA

```
[338]: # sns.pairplot(df_no_outliers2,hue="Price")
```

```
[376]: df_no_outliers2['Company'] = df_no_outliers2['Company'].astype('int')

numeric_columns = ['Weight(gm)', 'PPI', 'CPU_core', 'CPU_freq',
                    'Internal_mem(GB)', 'RAM', 'Dual_sim', 'RearCam', 'Front_Cam', 'Gen_5G',
                    'Battery', 'Thickness', 'Price']
df_no_outliers2[numeric_columns] = df_no_outliers2[numeric_columns].
    .astype('float')

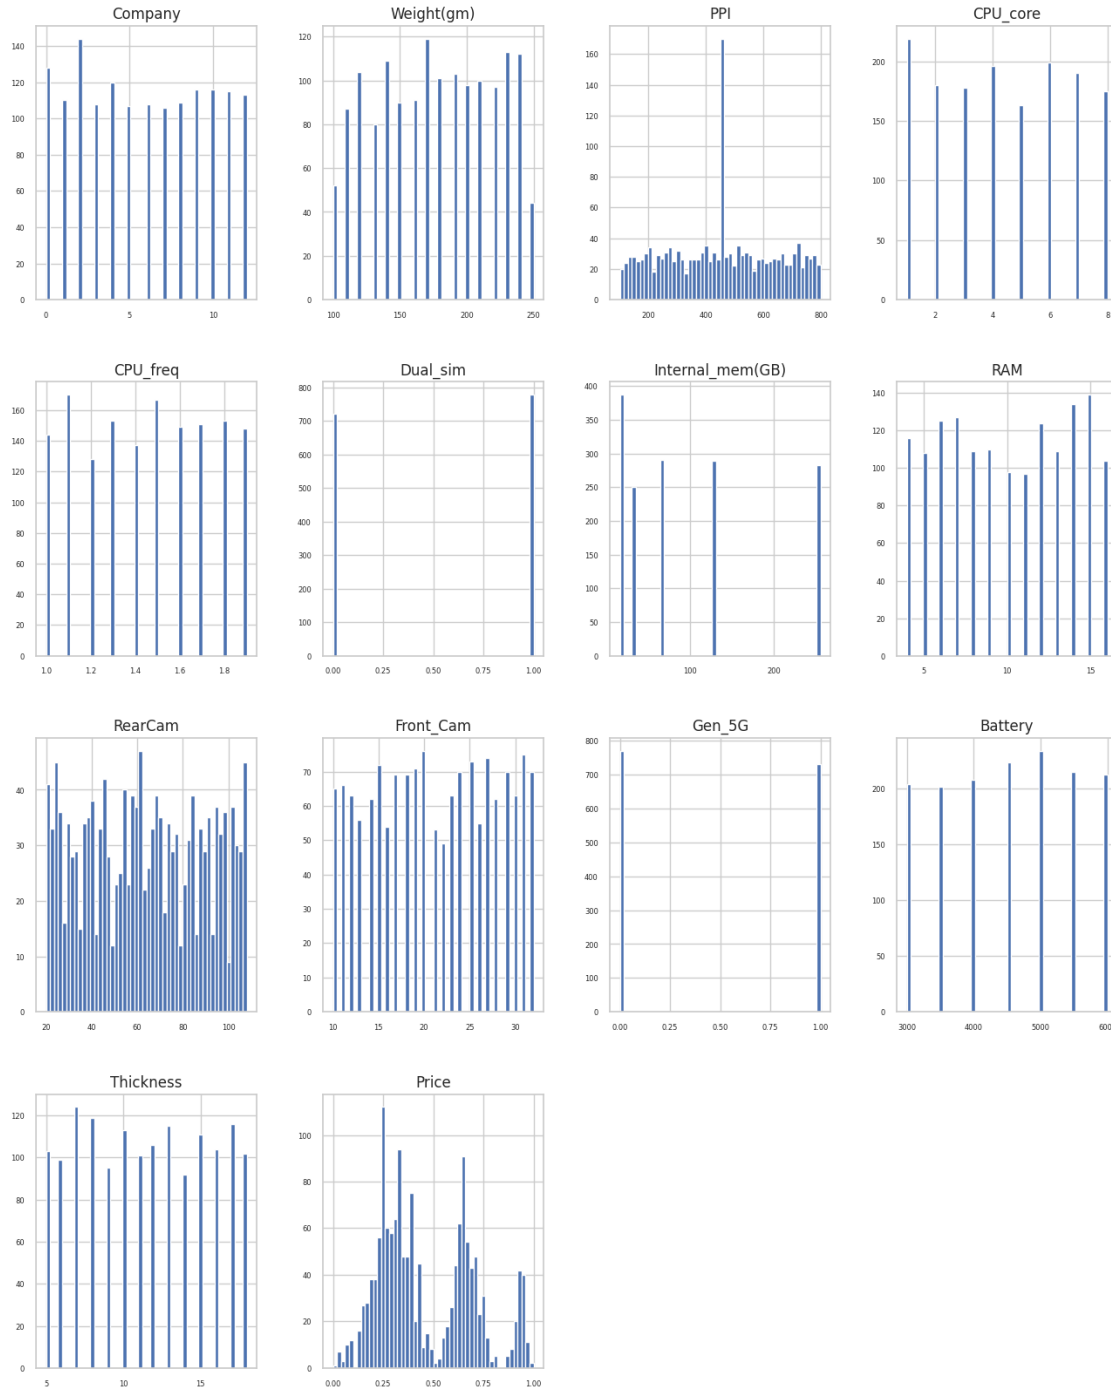
column_types = df_no_outliers2.dtypes
print(column_types)
```

```
Company          int64
```

Weight(gm)	float64
PPI	float64
CPU_core	float64
CPU_freq	float64
Dual_sim	float64
Internal_mem(GB)	float64
RAM	float64
RearCam	float64
Front_Cam	float64
Gen_5G	float64
Battery	float64
Thickness	float64
Price	float64
dtype:	object

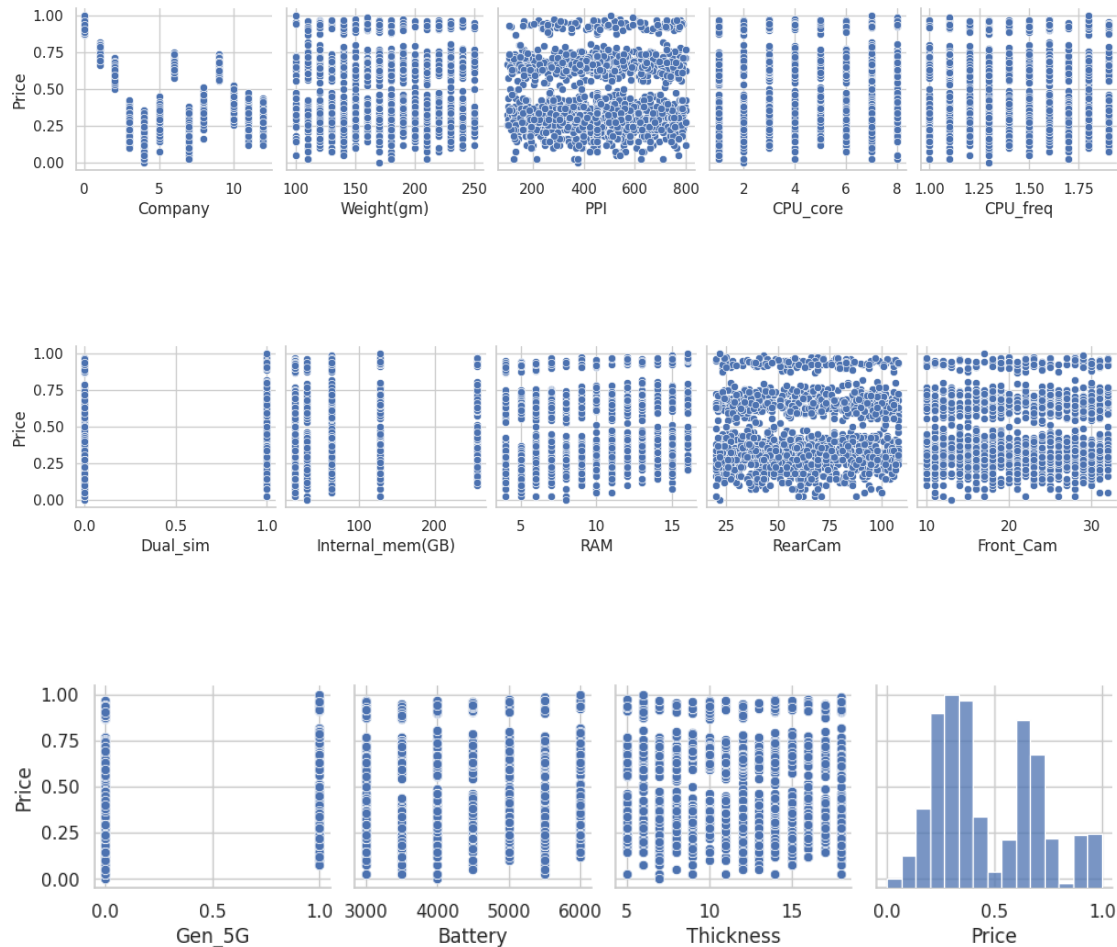
27 How mobile Prices are distribution across Companies

```
[380]: df_no_outliers2.hist(figsize=(16, 20), bins=50, xlabelsize=6, ylabelsize=6);
```




28 Pair Plots

```
[381]: for i in range(0, len(df_no_outliers2.columns), 5):
        sns.pairplot(data=df_no_outliers2,
                      x_vars=df_no_outliers2.columns[i:i+5],
                      y_vars=['Price'])
```




```
[340]: sns.lineplot(data=df_no_outliers2,x='Company',y='Price',color='blue')
plt.title('Company Vs Price',fontsize=20)
plt.xlabel('Company',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```



```
[341]: import matplotlib.pyplot as plt
company_encoding = {
    0: 'Apple',
    1: 'Google',
    2: 'Huawei',
    3: 'LG',
    4: 'Motorola',
    5: 'Nokia',
    6: 'OnePlus',
    7: 'Oppo',
    8: 'Realme',
    9: 'Samsung',
    10: 'Sony',
    11: 'Vivo',
    12: 'Xiaomi'
}
plt.figure(figsize=(10, 6))
plt.bar(df_no_outliers2['Company'], df_no_outliers2['Price'], color='skyblue')
```

```
plt.title('Company vs Price')
plt.xlabel('Company')
plt.ylabel('Price')
plt.xticks(df_no_outliers2['Company'], [company_encoding[company] for company_
    ↪ in df_no_outliers2['Company']], rotation=45, ha='right')

plt.show()
```




29 How the mobile price is distributed


```
[342]: print(df_no_outliers2['Price'].describe())
plt.figure(figsize=(9, 8))
sns.distplot(df_no_outliers2['Price'], color='g', bins=100, hist_kws={'alpha':
    ↪ 0.9});
```

```
count    1500.000000
mean      0.463338
std       0.236044
min       0.000000
25%       0.275045
50%       0.383050
```

```
75%          0.659710
max          1.000000
Name: Price, dtype: float64
```



```
[343]: plt.hist(df_no_outliers2['Price'], bins=20)
plt.title('Histogram of Price')
plt.show()
```



```
[345]: import matplotlib.pyplot as plt


# Set the figure size
plt.figure(figsize=(10, 6))

# Create a scatter plot
plt.scatter(df_no_outliers2['Internal_mem(GB)'], df_no_outliers2['Price'],
            color='green', alpha=0.5)

# Set title and labels
plt.title('Internal Memory vs Price')
plt.xlabel('Internal Memory (GB)')
plt.ylabel('Price')

# Set the x-axis labels based on the internal memory values
plt.xticks(df_no_outliers2['Internal_mem(GB)'].unique())

# Display the plot
plt.show()
```

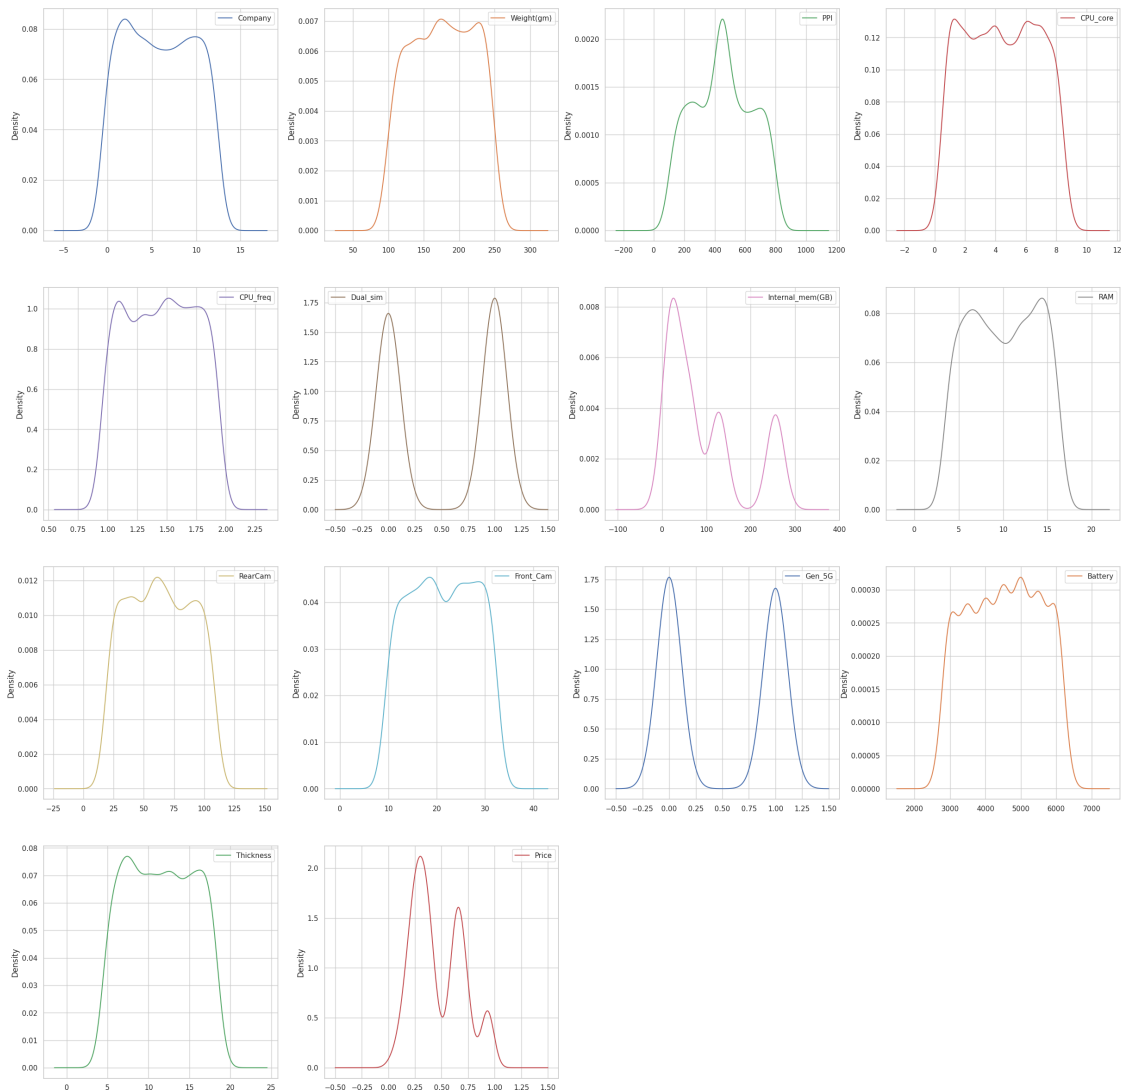


30 Plotting the distribution of quantitative features

```
[346]: import math
quantitative_features = df_no_outliers2.select_dtypes(include=['float64',
↪ 'int64']).columns
num_features = len(quantitative_features)

layout_rows = math.ceil(math.sqrt(num_features))
layout_cols = math.ceil(num_features / layout_rows)


df_no_outliers2[quantitative_features].plot(
    kind='density', subplots=True, layout=(layout_rows, layout_cols),
↪ sharex=False, figsize=(30, 30)
)
plt.show()
```



31 percentage of mobiles across the companies


```
[347]: import matplotlib.pyplot as plt
company_counts = df_no_outliers2['Company'].value_counts()
plt.figure(figsize=(10, 6))
plt.pie(company_counts, labels=company_counts.index, autopct='%1.1f%%',
        ↪startangle=90)
plt.title('Distribution of Mobiles by Company')
plt.show()
```

Distribution of Mobiles by Company




32 comparison of price & companies based on Ram

```
[348]: sns.lineplot(data=df_no_outliers2,x='Company',y='Price',hue='Internal_mem(GB)')
plt.title('Price Versus Company Based on RAM',fontsize=20)
plt.xlabel('Company',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```





33 Price Comparison by Company and 5G Generation

```
[349]: sns.barplot(data=df_no_outliers2, x='Company', y='Price', hue='Gen_5G')  
plt.title('Price Comparison by Company and 5G Generation')  
plt.show()
```



```
[350]: plt.figure(figsize=(10, 6))
sns.violinplot(x='RAM', y='CPU_freq', data=df_no_outliers2)
plt.xlabel('RAM')
plt.ylabel('CPU_freq')
plt.title('RAM vs CPU_freq')
plt.show()
```



```
[351]: df_no_outliers2.head()
```

```
[351]:
```

	Company	Weight(gm)	PPI	CPU_core	CPU_freq	Dual_sim	Internal_mem(GB)	\
0	9	180.0	312	1	1.5	0	16.0	
1	2	160.0	362	1	1.6	0	64.0	
2	1	160.0	241	1	1.0	0	16.0	
3	6	210.0	555	4	1.2	1	32.0	
4	11	100.0	607	8	1.8	1	128.0	

	RAM	RearCam	Front_Cam	Gen_5G	Battery	Thickness	Price
0	8	77	31	1.0	5000.0	17	0.594341
1	8	51	13	0.0	3000.0	12	0.520853
2	11	84	23	0.0	5000.0	11	0.685658
3	7	91	10	0.0	6000.0	8	0.659710
4	12	71	27	0.0	4500.0	11	0.368775

34 Price Comparison based on Dual sim feature

```
[352]: import plotly.express as px

# Assuming df_no_outliers1 is your DataFrame
fig = px.scatter(df_no_outliers2, x="Dual_sim", y="Price")
```

```

# Define custom y-axis tick values and labels
custom_x_ticks = [0, 1]

# Update layout to set custom y-axis ticks
fig.update_layout(
    xaxis=dict(
        tickmode='array',
        tickvals=custom_x_ticks,
    ),
    width=800,
    height=600,
)
fig.show()

```


35 comparison of price & Ram

```

[353]: import matplotlib.pyplot as plt
import seaborn as sns

# Assuming df is your DataFrame
plt.figure(figsize=(10, 6))
sns.scatterplot(x='RAM', y='Price', data=df_no_outliers2)
plt.title('Comparison between RAM and Price')
plt.xlabel('RAM')
plt.ylabel('Price')
plt.grid(True)
plt.show()

```




```
[354]: from matplotlib import cm

cl = cm.get_cmap('Accent')
tw1 = df_no_outliers2.query('Gen_5G==1 and Dual_sim==1')['Price'].value_counts()
tw2 = df_no_outliers2.query('Gen_5G==0 and Dual_sim==1')['Price'].value_counts()
tw3 = df_no_outliers2.query('Gen_5G==0 and Dual_sim==0')['Price'].value_counts()
tw4 = df_no_outliers2.query('Gen_5G==1 and Dual_sim==0')['Price'].value_counts()

colors = [cl(0), cl(0.2), cl(0.9), cl(0.4)]
twlist = [tw1, tw2, tw3, tw4]
titles = ['Gen_5G==1 and Dual_sim==1', 'Gen_5G==0 and Dual_sim==1', 'Gen_5G==0 and Dual_sim==0', 'Gen_5G==1 and Dual_sim==0']


fig, axis = plt.subplots(4, 1, figsize=(30, 20))
for i in range(4):
    sns.barplot(y=twlist[i], x=twlist[i].index, ax=axis[i], color=colors[i])
    axis[i].set_title(titles[i])
    axis[i].set_xticks(range(len(twlist[i].index)))
    axis[i].set_xticklabels(twlist[i].index, rotation=45, ha='right')

fig.tight_layout()
plt.show()
```




36 comparison of price & Battery

```
[355]: sns.swarmplot(data=df_no_outliers2, x='Battery', y='Price')
plt.title('Battery Vs Price')
plt.show()
```



37 comparison of price & cpu_core


```
[356]: d_grouped = df_no_outliers2.groupby('CPU_core')[['CPU_core', 'Price']].max()
plt.stem(d_grouped['CPU_core'], d_grouped['Price'])
plt.title('Maximum Price by CPU Core')
plt.show()
```



```
[357]: pip install squarify
```


```
Collecting squarify
  Downloading squarify-0.4.3-py3-none-any.whl (4.3 kB)
Installing collected packages: squarify
Successfully installed squarify-0.4.3
```

```
[358]: import squarify
d1_grouped = df_no_outliers2[['CPU_core', 'Price']].groupby('CPU_core').mean().
    ↪reset_index()
# Creating a treemap
squarify.plot(sizes=d1_grouped['Price'], label=d1_grouped['CPU_core'])
plt.title('Average mobile price by cpu core')
plt.show()
```


38 comparison of price & cpu freq

```
[359]: import squarify
d1_grouped = df_no_outliers2[['CPU_freq', 'Price']].groupby('CPU_freq').mean().
    ↪reset_index()
# Creating a treemap
squarify.plot(sizes=d1_grouped['Price'], label=[f'{price:.2f}' for freq, price_
    ↪in zip(d1_grouped['CPU_freq'], d1_grouped['Price'])])
plt.title('Average mobile price by CPU_freq')
plt.show()
```



```
[360]: df_no_outliers2['PPI'].plot(kind='hist',bins=10, color='pink')
```


```
[360]: <Axes: ylabel='Frequency'>
```



39 comparison of Ram & Battery & Company

```
[361]: sns.relplot(x='Company', y='RAM', hue='Battery', data=df_no_outliers2)
```

```
[361]: <seaborn.axisgrid.FacetGrid at 0x78fec30ec490>
```



40 PPI Vs Company

```
[362]: sns.lineplot(data=df_no_outliers2,x='PPI',y='Company',color='blue')
plt.title('PPI Vs Company',fontsize=20)
plt.xlabel('PPI',fontsize=15)
plt.ylabel('Company',fontsize=15)
plt.show()
```

